

Zhipeng Tian  
P436  
Assignment 3B  
3/31/12

**1. List the kernel functions that are called inside the system call sys mmap.**

sys\_mmap()---->

do\_mmap2()---->

do\_mmap\_pgoff()---->

get\_unmapped\_area()  
vma\_merge()  
deny\_write\_access()  
shmem\_zero\_setup()  
find\_vma\_prepare()  
vma\_link()  
make\_pages\_present()  
zap\_page\_range()

**2. List the synchronization primitives and variables that the kernel uses when manipulating the memory-related data structures, mm struct and vm area struct.**

**mm\_struct:**

struct vm\_area\_struct \* mmap;  
struct rb\_root mm\_rb;  
struct vm\_area\_struct \* mmap\_cache;  
unsigned long mmap\_base;  
unsigned long task\_size;  
unsigned long cached\_hole\_size;  
unsigned long free\_area\_cache;  
pgd\_t \* pgd;  
atomic\_t mm\_users;  
atomic\_t mm\_count;  
int map\_count;  
spinlock\_t page\_table\_lock;  
struct rw\_semaphore mmap\_sem;  
struct list\_head mmlist;  
unsigned long hiwater\_rss;  
unsigned long hiwater\_vm;  
unsigned long total\_vm, locked\_vm, shared\_vm, exec\_vm;  
unsigned long stack\_vm, reserved\_vm, def\_flags, nr\_ptes;  
unsigned long start\_code, end\_code, start\_data, end\_data;  
unsigned long start\_brk, brk, start\_stack;  
unsigned long arg\_start, arg\_end, env\_start, env\_end;

```

unsigned long saved_auxv[AT_VECTOR_SIZE];
struct mm_rss_stat rss_stat;
struct linux_binfmt *binfmt;
cpumask_var_t cpu_vm_mask_var;
mm_context_t context;
unsigned int faultstamp;
unsigned int token_priority;
unsigned int last_interval;
atomic_t oom_disable_count;
unsigned long flags;
struct core_state *core_state;
spinlock_t ioctx_lock;
struct hlist_head ioctx_list;
struct task_struct __rcu *owner;
struct file *exe_file;
unsigned long num_exe_file_vmas;
struct mmu_notifier_mm *mmu_notifier_mm;
pgtable_t pmd_huge_pte; /* protected by page_table_lock */
struct cpumask cpumask_allocation;

```

#### **vm\_area\_struct:**

```

struct mm_struct * vm_mm;
unsigned long vm_start;
unsigned long vm_end;
struct vm_area_struct *vm_next, *vm_prev;
pgprot_t vm_page_prot;
unsigned long vm_flags;
struct rb_node vm_rb;
union {
    struct {
        struct list_head list;
        void *parent; /* aligns with prio_tree_node parent */
        struct vm_area_struct *head;
    } vm_set;
    struct raw_prio_tree_node prio_tree_node;
} shared;
struct list_head anon_vma_chain;
struct anon_vma *anon_vma;
const struct vm_operations_struct *vm_ops;
unsigned long vm_pgoff;
struct file * vm_file;
void * vm_private_data;
struct vm_region *vm_region;
struct mempolicy *vm_policy;

```

### 3. What does the function find vma do and how does it work?

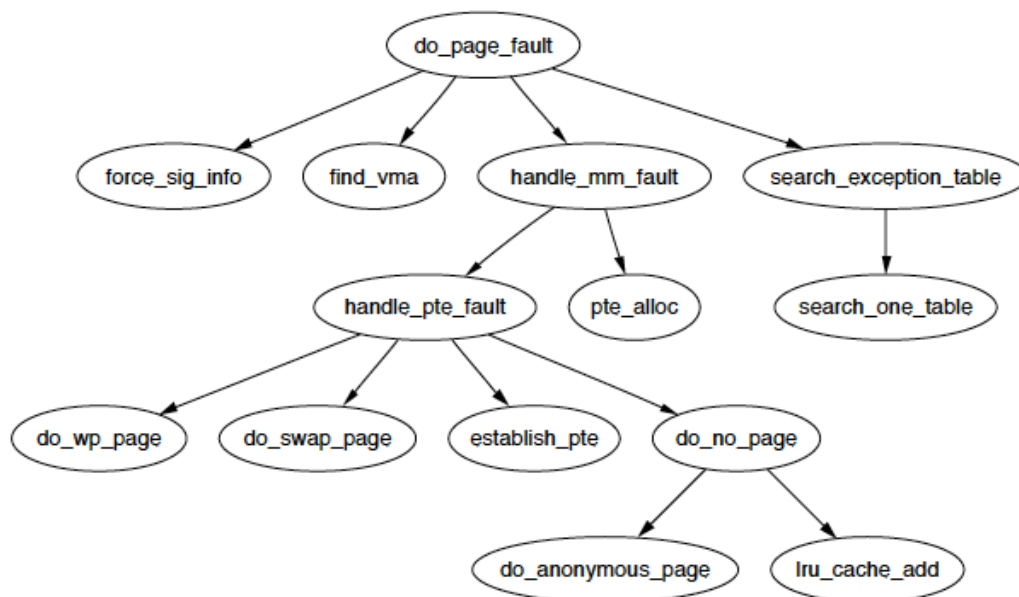
The `find_vma()` function locates the first memory region which satisfies `addr < vm_end` and returns the address of its descriptor, returns a NULL pointer if nothing is found.

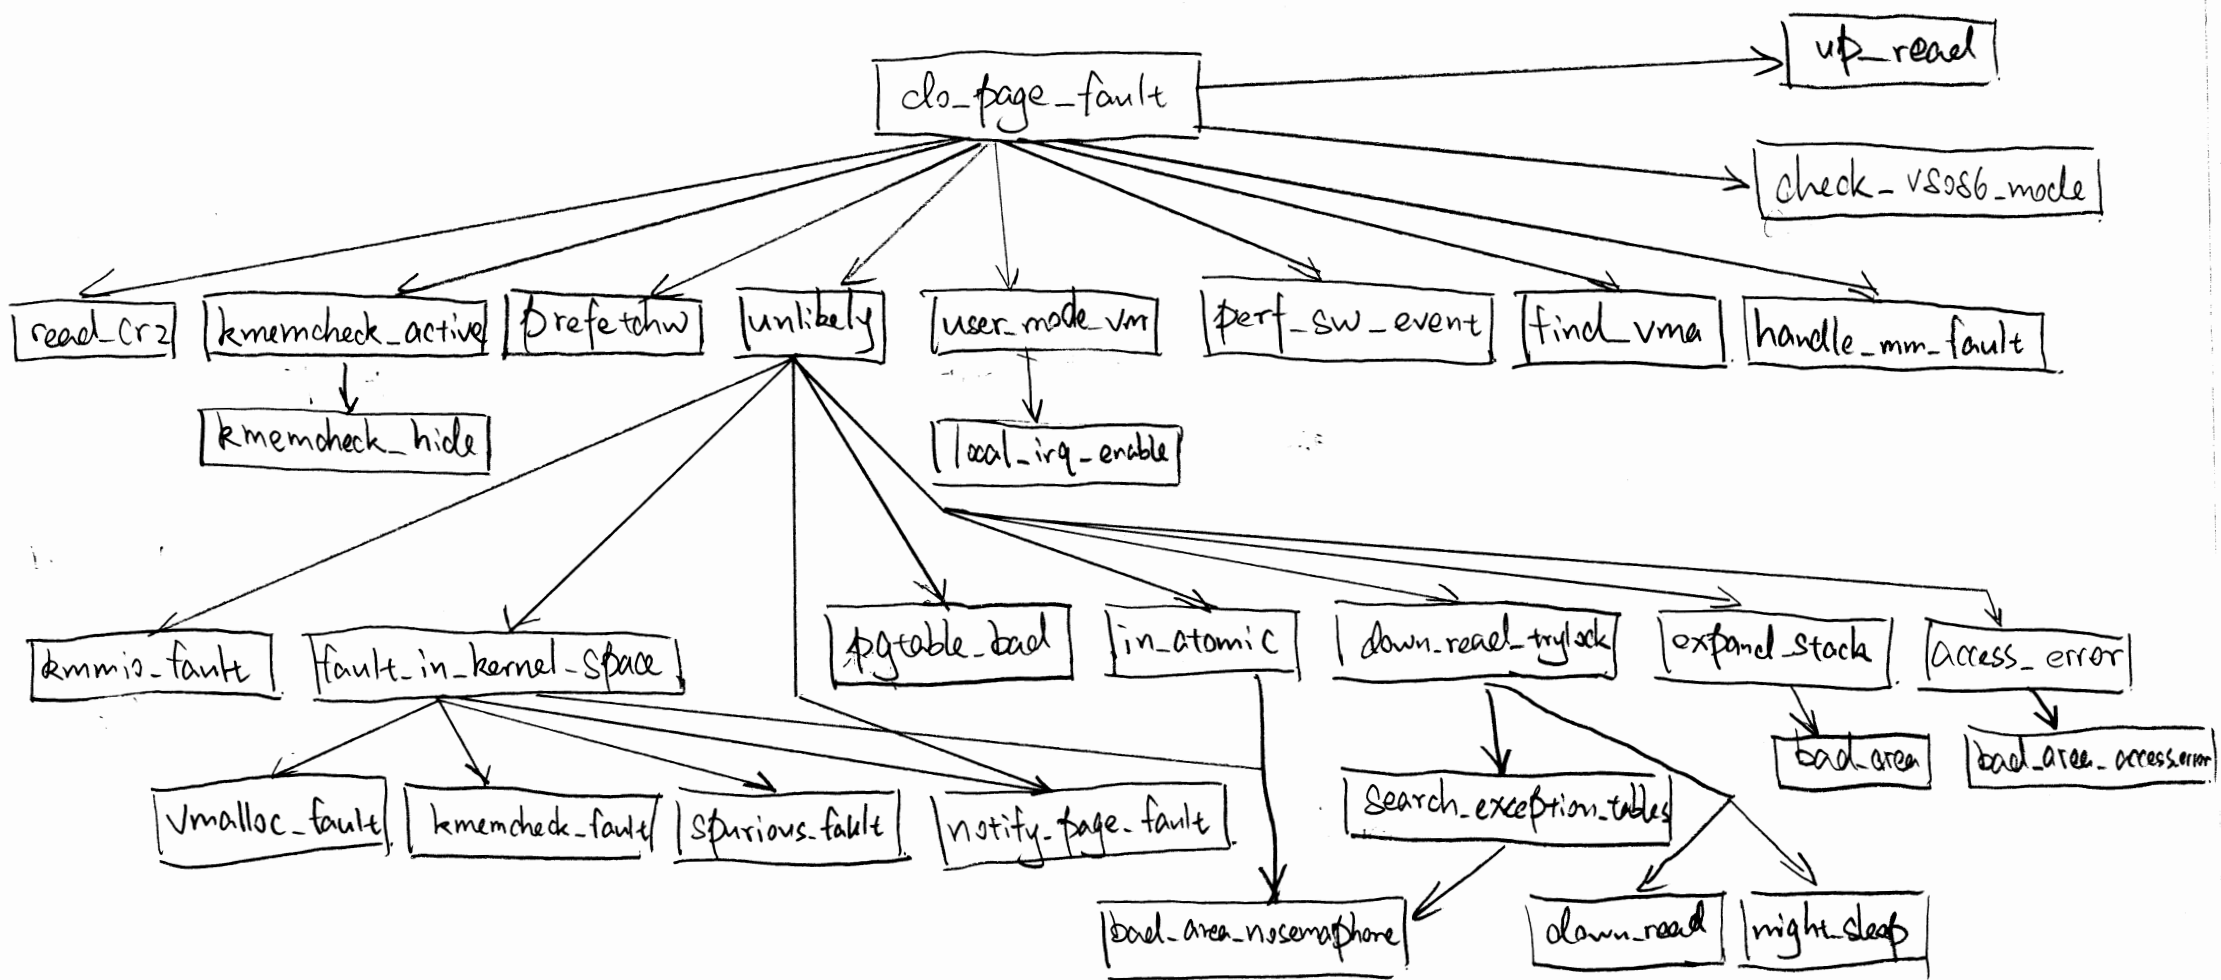
It checks cache first, checks whether the region identified by `mmap_cache` has `addr` inside it. If so, it returns the region descriptor pointer. Otherwise, it will scan the memory region of the process by looking up the memory region in the red-black tree.

**4. A call graph is a graph in which vertices are function names and there is an edge from a vertex `f1` to `f2` if function `f1` calls function `f2`. Notice that a call graph may have self edges (due to recursion) and cycles (due to mutual recursion). Draw a call graph for `do_page_fault`. This should include all the helper functions called under different circumstances.**

The first call graph is the online resource I found, [my own version of `do\_page\_fault\(\)` call graph is on the next page](#), and it is based on the file `arch/x86/mm/fault.c`.

Here is the call graph created by **Mel Gorman** in his “**Code Commentary On The Linux Virtual Memory Manager**”. It shows the call graph on a deeper level than me, which helped me understand the `do_page_fault()` much better.





\* arch/x86/mm/fault.c

\* All functions calls are inside do\_page\_fault().

Zhifeng Tian