

Zhipeng Tian
P436
Assignment 3B
3/31/12

1. List the kernel functions that are called inside the system call sys mmap.

sys_mmap()---->

do_mmap2()---->

do_mmap_pgoff()---->

get_unmapped_area()
vma_merge()
deny_write_access()
shmem_zero_setup()
find_vma_prepare()
vma_link()
make_pages_present()
zap_page_range()

2. List the synchronization primitives and variables that the kernel uses when manipulating the memory-related data structures, mm struct and vm area struct.

mm_struct:

struct vm_area_struct * mmap;
struct rb_root mm_rb;
struct vm_area_struct * mmap_cache;
unsigned long mmap_base;
unsigned long task_size;
unsigned long cached_hole_size;
unsigned long free_area_cache;
pgd_t * pgd;
atomic_t mm_users;
atomic_t mm_count;
int map_count;
spinlock_t page_table_lock;
struct rw_semaphore mmap_sem;
struct list_head mmlist;
unsigned long hiwater_rss;
unsigned long hiwater_vm;
unsigned long total_vm, locked_vm, shared_vm, exec_vm;
unsigned long stack_vm, reserved_vm, def_flags, nr_ptes;
unsigned long start_code, end_code, start_data, end_data;
unsigned long start_brk, brk, start_stack;
unsigned long arg_start, arg_end, env_start, env_end;

```

unsigned long saved_auxv[AT_VECTOR_SIZE];
struct mm_rss_stat rss_stat;
struct linux_binfmt *binfmt;
cpumask_var_t cpu_vm_mask_var;
mm_context_t context;
unsigned int faultstamp;
unsigned int token_priority;
unsigned int last_interval;
atomic_t oom_disable_count;
unsigned long flags;
struct core_state *core_state;
spinlock_t ioctx_lock;
struct hlist_head ioctx_list;
struct task_struct __rcu *owner;
struct file *exe_file;
unsigned long num_exe_file_vmas;
struct mmu_notifier_mm *mmu_notifier_mm;
pgtable_t pmd_huge_pte; /* protected by page_table_lock */
struct cpumask cpumask_allocation;

```

vm_area_struct:

```

struct mm_struct * vm_mm;
unsigned long vm_start;
unsigned long vm_end;
struct vm_area_struct *vm_next, *vm_prev;
pgprot_t vm_page_prot;
unsigned long vm_flags;
struct rb_node vm_rb;
union {
    struct {
        struct list_head list;
        void *parent; /* aligns with prio_tree_node parent */
        struct vm_area_struct *head;
    } vm_set;
    struct raw_prio_tree_node prio_tree_node;
} shared;
struct list_head anon_vma_chain;
struct anon_vma *anon_vma;
const struct vm_operations_struct *vm_ops;
unsigned long vm_pgoff;
struct file * vm_file;
void * vm_private_data;
struct vm_region *vm_region;
struct mempolicy *vm_policy;

```

3. What does the function find vma do and how does it work?

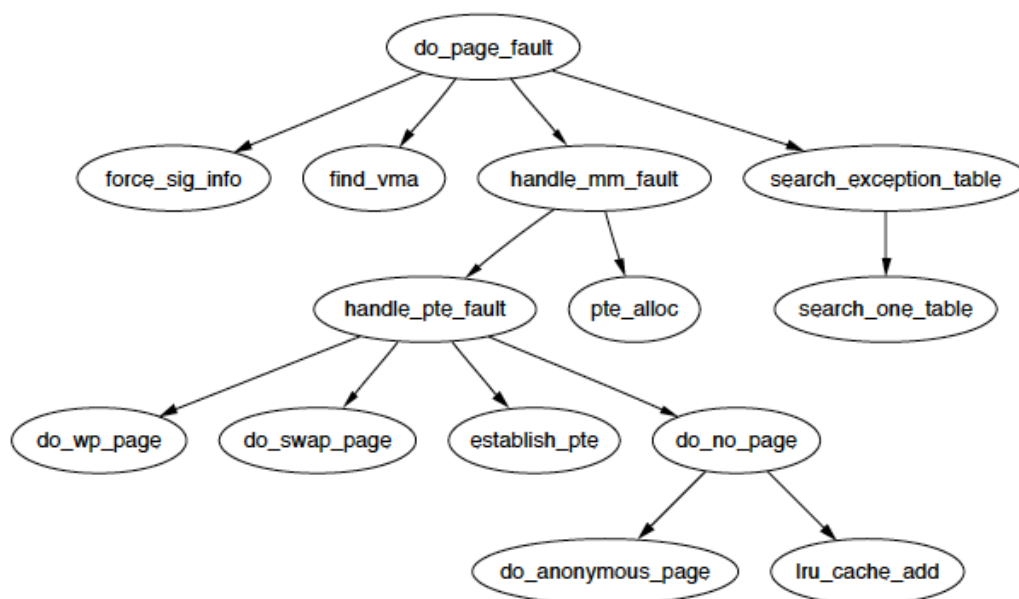
The `find_vma()` function locates the first memory region which satisfies `addr < vm_end` and returns the address of its descriptor, returns a NULL pointer if nothing is found.

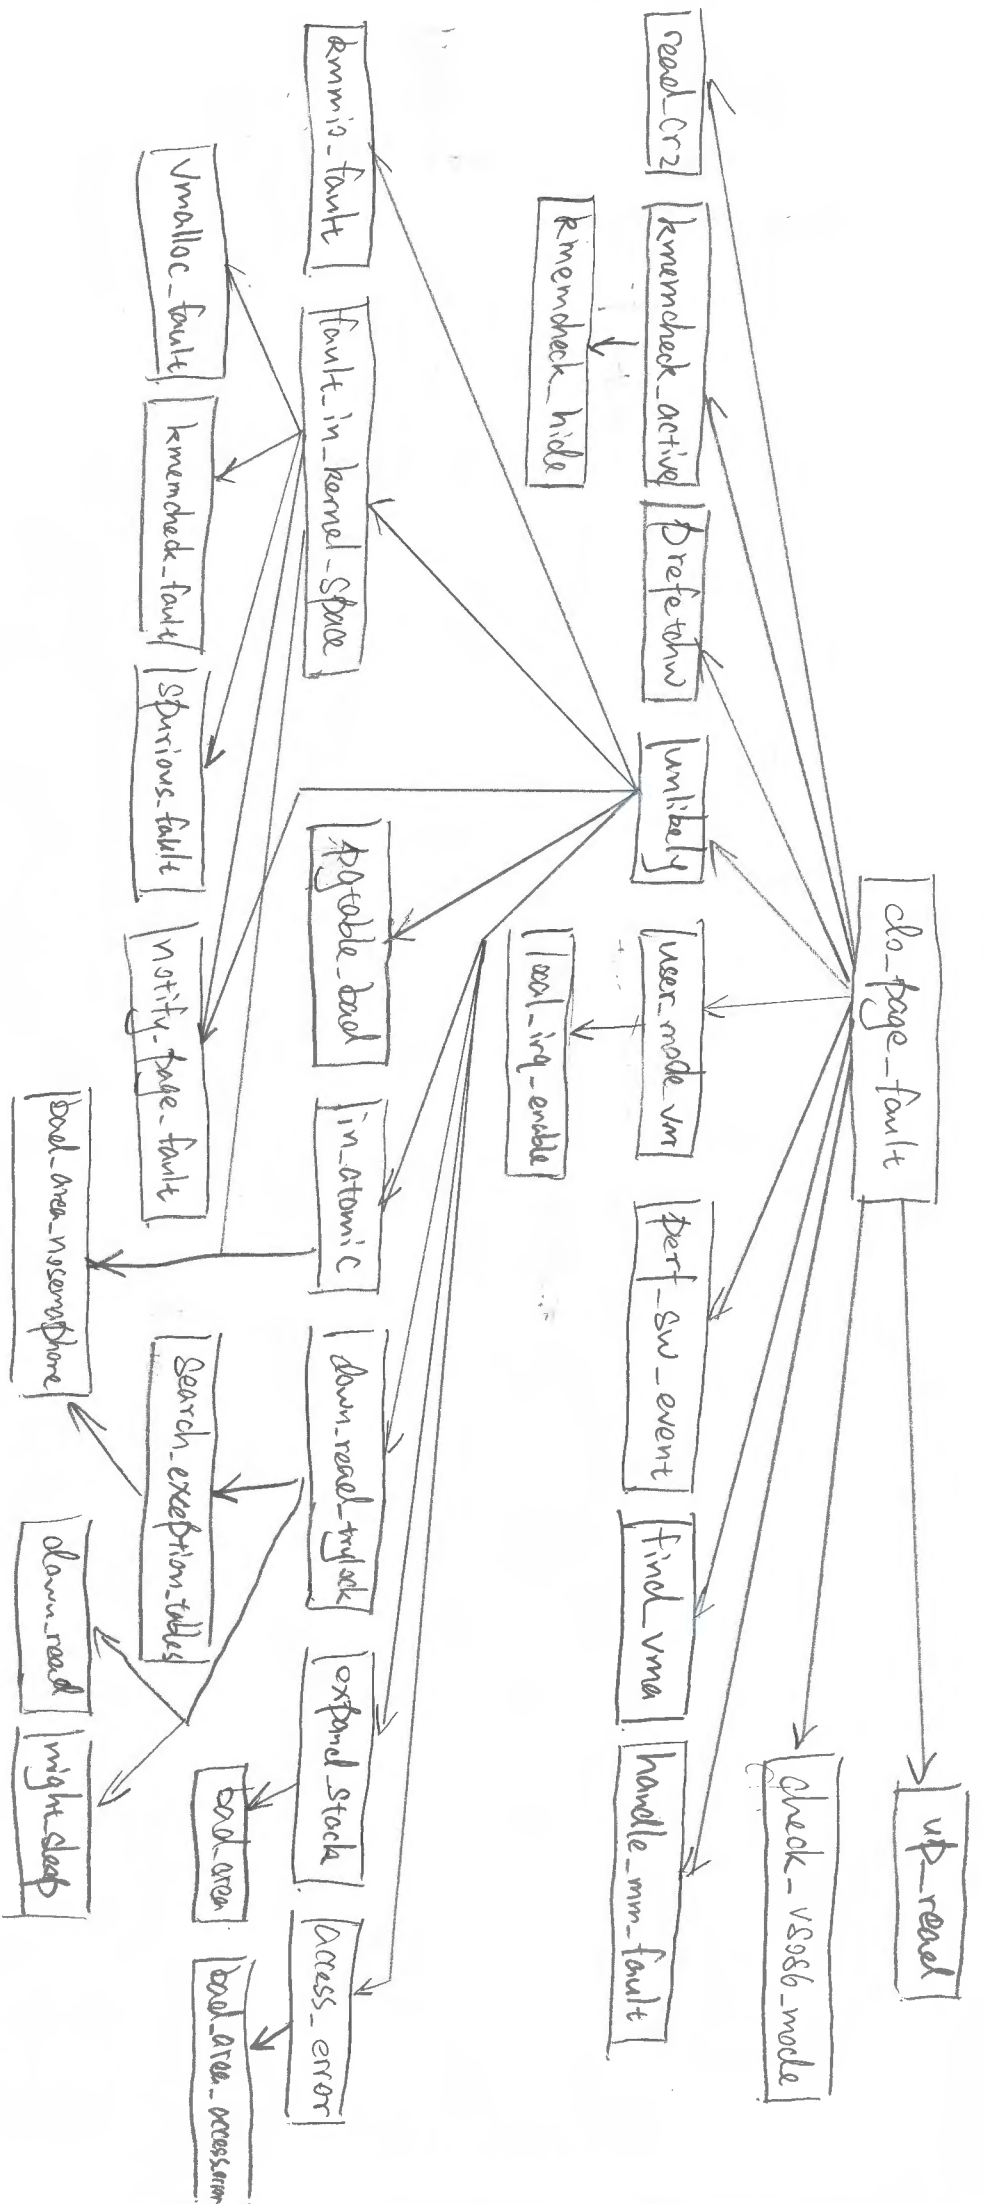
It checks cache first, checks whether the region identified by `mmap_cache` has `addr` inside it. If so, it returns the region descriptor pointer. Otherwise, it will scan the memory region of the process by looking up the memory region in the red-black tree.

4. A call graph is a graph in which vertices are function names and there is an edge from a vertex f1 to f2 if function f1 calls function f2. Notice that a call graph may have self edges (due to recursion) and cycles (due to mutual recursion). Draw a call graph for do page fault. This should include all the helper functions called under different circumstances.

The first call graph is the online resource I found, [my own version of do_page_fault\(\) call graph is on the next page](#), and it is based on the file `arch/x86/mm/fault.c`.

Here is the call graph created by **Mel Gorman** in his “**Code Commentary On The Linux Virtual Memory Manager**”. It shows the call graph on a deeper level than me, which helped me understand the `do_page_fault()` much better.





- * arch/x86/mm/fault.c
- * All functions calls are inside do_page_fault()

ZhiFeng Tan